

Launching and Scaling Enterprise AI Enablement, Literacy, and Pilot Services

Lessons from large corporations, with emphasis on controlled use of corporate information

Comparative research report | August 2026

Scope note. This report synthesizes official corporate publications, annual reports, public company materials, and clearly identified vendor case studies. Public information does not reveal all internal controls. The report distinguishes published facts, cross-company observations, and recommendations. It is not legal, privacy, security, employment, procurement, regulatory, or financial advice.

Executive summary

Large corporations are generally not scaling generative AI by giving employees unrestricted access to public tools. The recurring model is a controlled enterprise environment paired with central governance, federated business ownership, role-based learning, bounded pilots, technical controls, and evidence-based progression to production.

The strongest public examples show different entry points. JPMorganChase built a shared controlled platform and made it available to more than 200,000 employees. Morgan Stanley began with a narrowly scoped, knowledge-grounded assistant for financial advisers and used formal evaluations before expansion. Walmart first created a safe internal playground, then launched a proprietary associate assistant and expanded it internationally. Moderna used an internal chatbot to build familiarity, tested competing enterprise tools, pursued broad proficiency, and enabled employees to create function-specific assistants. Unilever’s public materials indicate a formal governance platform covering use-case registration, risk triage, assurance evidence, lifecycle monitoring, and cross-functional oversight. IBM publicly describes end-to-end use-case inventories, factsheets, approval workflows, monitoring, and cross-platform governance.

The most important information-protection conclusion is that contractual non-training terms are necessary but insufficient. Enterprises also need institutional identity, data authority, classification, DLP and retention settings, repository permission review, connector restrictions, secrets management, human accountability, monitoring, incident response, and periodic revalidation.

A practical operating model is a central AI enablement and governance service that supplies approved platforms, training, intake, reusable patterns, and portfolio visibility, while business units remain accountable for purpose, data authority, workflow ownership, human review, and outcomes. Higher-risk integrations and agentic actions should receive substantially more review than protected chat or synthetic-data learning exercises.

1. Scope and method

The report examines how selected large corporations publicly describe the introduction and scaling of enterprise generative AI services. The sample was selected to represent regulated financial services, retail, life sciences, consumer goods, and enterprise technology. It is not a ranking and is not statistically representative.

Evidence category	Meaning	Treatment in this report
Company-published fact	A statement in an annual report, corporate press release, policy, or official company page.	Used as direct evidence, while recognizing that public descriptions may be selective.
Vendor case study	A customer story published by a technology provider, usually with customer participation.	Used with an explicit source label; performance and adoption claims are not treated as independently audited.
Cross-company observation	A recurring pattern supported by several sources.	Presented as synthesis rather than as a fact about every company.
Recommendation	A proposed practice derived from the comparison.	Requires adaptation to industry, jurisdiction, risk appetite, architecture, and corporate policy.

2. Comparative corporate cases

2.1 JPMorganChase: shared controlled platform with federated business use

JPMorganChase reported that it launched LLM Suite in 2024 to more than 200,000 employees worldwide. The firm described it as a controlled desktop environment that protects customer and company data and provides shared capabilities for business teams to build and integrate generative-AI use cases. Its public materials also describe a centralized Chief Data & Analytics Office that sets firmwide strategy, governance, risk controls, and tooling, while embedded teams in the businesses deliver use cases with approved tools. [C-1][C-2]

Operational implication: a large enterprise can separate the platform and governance layer from business delivery. The central function establishes common controls and infrastructure; business teams remain responsible for workflow value and implementation.

2.2 Morgan Stanley: narrow knowledge use case, formal evaluation, then expansion

Morgan Stanley's Firmwide AI team publicly states that it unifies AI efforts, protects data and systems, and incorporates controls and oversight. Its first widely publicized wealth-management assistant grounded responses in the firm's intellectual capital. The company later expanded to meeting summarization, with adviser review before finalization. Morgan Stanley reported 98% adoption among financial-adviser teams. [C-3][C-4]

A vendor case study describes an evaluation framework in which expert reviewers graded accuracy and coherence, retrieval methods were refined, daily regression testing was used, and separate evaluation datasets were created for meeting types. The same case study states that zero data retention was part of the vendor arrangement. [V-1]

Operational implication: a credible path to scale is to begin with a bounded, high-value knowledge problem, build domain-specific evaluation datasets, require human review, and expand only after reliability and adoption evidence are established.

2.3 Walmart: internal playground, proprietary assistant, employee feedback, global expansion

Walmart Global Tech states that the company first provided a GenAI Playground as a safe space for associates to experiment, then launched My Assistant for drafting, summarization, and ideation. Walmart reported that the first version moved from concept to a viable product in about 60 days and was initially introduced to 50,000 U.S. campus associates before international expansion. The company also states that associate feedback was incorporated during rollout. [C-5][C-6]

Operational implication: broad learning access and a simple internal assistant can precede more specialized developer or workflow tools. A safe experimentation environment can reduce unmanaged use while generating demand and feedback.

2.4 Moderna: broad proficiency target, user testing, and employee-built assistants

An OpenAI customer case study reports that Moderna first launched an internal chatbot, mChat, which was adopted by more than 80% of employees. Moderna then compared mChat, Copilot, and ChatGPT

Enterprise through user testing before selecting a preferred enterprise platform. The company set an objective of broad adoption and proficiency and enabled employees to create function-specific GPTs. The case study reports 750 GPTs within two months, with examples in clinical analysis, contracts, internal policy retrieval, and communications. [V-2]

The same case study describes a clinical analysis assistant as supporting, rather than replacing, the study team’s judgment and states that the tool provides rationale, references, and visualizations for human review. [V-2]

Operational implication: adoption can be treated as a workforce capability program, but rapid creation of persistent assistants requires ownership, source approval, lifecycle controls, and an inventory to prevent unmanaged proliferation.

2.5 Unilever: operationalizing governance as a product and workflow

Current Unilever role descriptions provide unusually specific public evidence about its governance operating model. The company describes a governance platform supporting AI use-case registration, risk triage, assessment, assurance workflows, evidence capture, lifecycle monitoring, adoption metrics, and collaboration among technology, legal, cyber, privacy, finance, procurement, and business teams. [C-7]

Other public role materials describe moving initiatives from discovery through proof of concept, pilot, and scale-up, and identify AI councils, governance frameworks, oversight, security, data protection, training, and centers of excellence. [C-8][C-9]

Operational implication: governance is more scalable when implemented as a service and product workflow, not solely as policy documents or committee review. Intake quality, structured evidence, workflow status, and adoption metrics become operational assets.

2.6 IBM: inventory, factsheets, lifecycle workflow, and continuous monitoring

IBM’s public documentation describes an AI governance process that begins with a business user creating an AI use case, associates models or prompt templates with the use case, records development and evaluation details in factsheets, enables validator and risk-officer review, and monitors deployed assets. The platform supports third-party models as well as IBM models. [C-10][C-11]

IBM also reports using its governance platform internally to approve more than 1,000 models for reuse and reduce data-clearance processing time, although this is a vendor-reported product claim rather than an independent audit. [C-12]


Operational implication: an enterprise inventory should connect purpose, owners, models, prompts, datasets, evaluations, risks, controls, deployment status, and production monitoring rather than functioning as a static register.

2.7 Comparative findings

Company	Published starting pattern	Control emphasis	Scaling mechanism
JPMorganChase	Firmwide controlled platform	Central strategy, governance, risk controls, protected data environment	Shared capabilities plus embedded business teams

Company	Published starting pattern	Control emphasis	Scaling mechanism
Morgan Stanley	Narrow knowledge assistant	Expert evaluation, regression testing, human review, controlled vendor arrangement	Expand after reliability and adoption evidence
Walmart	Safe playground followed by proprietary assistant	Internal environment, familiar use cases, feedback	Initial large cohort, then international rollout
Moderna	Internal chatbot, tool comparison, broad workforce access	Enterprise platform, human-supported clinical judgment	Employee-created function-specific assistants
Unilever	Governance platform and innovation funnel	Intake, triage, assurance evidence, lifecycle monitoring	Cross-functional governance embedded in delivery
IBM	Use-case inventory and lifecycle governance	Factsheets, approval workflow, evaluation and monitoring	Reusable governed assets across platforms

3. Enterprise control architecture



A vendor non-training commitment is important, but it is only one control layer.

Figure 1. Recommended control layers for enterprise AI services.

3.1 Contract and vendor controls

- Define whether prompts, files, outputs, feedback, telemetry, and metadata may be used for model training or service improvement.
- Confirm retention, deletion, support access, subprocessors, data location, breach notification, audit rights, intellectual property, indemnity, termination, and data export.
- Validate whether commitments differ by model, feature, connector, API, region, licence tier, or administrator setting.
- Require notice and reassessment for material product, model, contract, or data-handling changes.

3.2 Corporate identity and access

- Require enterprise-managed accounts, single sign-on, multifactor authentication, and automated joiner/mover/leaver processes.
- Use role-based groups, workspaces, and pilot cohorts rather than unmanaged personal accounts.
- Review administrator, vendor-support, delegated, guest, export, and conversation-sharing privileges.
- Use conditional access, device controls, and geographic restrictions where appropriate.

3.3 Data authority, classification, and minimization

- Separate approval of a platform from authority to use a particular dataset for a particular purpose.
- Identify data owners and restrictions arising from customer contracts, employee privacy, intellectual property, confidentiality, regulation, litigation holds, export controls, and records requirements.
- Prefer public, synthetic, de-identified, or reduced data for learning and feasibility testing.
- Document permitted and prohibited information for each service, feature, connector, and pilot.

3.4 DLP, retention, records, and monitoring

- Apply data-loss-prevention policies and sensitivity labels where supported; test whether they work across prompts, uploads, generated outputs, sharing, and connectors.
- Define retention for chats, uploaded files, agent knowledge, logs, embeddings, indexes, and exports.
- Clarify when prompts and outputs become corporate records and how legal holds, investigation, and e-discovery requirements apply.
- Monitor adoption, policy exceptions, anomalous activity, shadow AI, incidents, model changes, and control effectiveness.

4. Classify the interaction pattern, not only the product

Interaction pattern determines control intensity		
1. Protected chat	Drafting, summarizing, brainstorming	Low persistence and integration; authorization and output handling still matter.
2. Uploaded files	Document analysis in an enterprise workspace	Retention, metadata, sharing, and records obligations increase.
3. Persistent assistants	Custom instructions, shared knowledge	Ownership, stale content, access recertification, and decommissioning are required.
4. Connected retrieval / APIs	PLG, connectors, workflow integration	Oversharing, secrets, logs, subprocessors, and software lifecycle risks increase.
5. Agentic actions	Email, tickets, records, transactions, decisions	Errors can propagate into business systems and affect customers or employees.

Figure 2. Control intensity should increase with persistence, integration, autonomy, and potential impact.

Interaction pattern	Primary additional risks	Minimum additional controls
Protected chat	Unauthorized disclosure, inaccurate output, inappropriate reuse	Approved account, data guidance, user training, human verification, retention settings
Uploaded-file analysis	Persistent copies, hidden metadata, broad sharing, records and IP obligations	File/data inventory, upload limits, retention/deletion, sharing controls, approved purpose
Persistent assistant	Stale knowledge, unclear owner, unmanaged sharing, instruction drift	Named owner, approved sources, versioning, access review, usage logging, decommission date
Repository-connected retrieval	Overshared repositories, stale groups, cross-document inference, prompt injection	Permission audit, security trimming test, source citations, adversarial tests, index deletion process
API workflow	Secrets, logs, retries, downstream data stores, code vulnerabilities	Architecture/security review, secrets vault, logging standard, testing, change control, incident plan
Agentic action	Unauthorized or erroneous business action, cascading errors, customer/employee impact	Least privilege, approval gates, transaction limits, rollback, high-quality audit trail, continuous monitoring

5. Recommended enterprise service model

Cross-company evidence supports a hub-and-spoke model. A central service provides common platforms, governance workflows, training, standards, evaluation methods, and portfolio visibility. Business, functional, or regional teams own the problem, process, authorized data, user adoption, human review, and measurable outcomes.

Service component	Purpose	Typical outputs
AI advisory front door	Turn questions and ideas into clear next steps	Triage decision, consultation, referral, use-case framing
Enterprise learning environment	Provide a safe alternative to unmanaged public-tool use	Protected chat/playground, sample tasks, synthetic datasets
Role-based literacy and enablement	Build practical capability for users, managers, developers, data owners, and reviewers	Training paths, office hours, champions network, scenario exercises
Use-case intake and risk tiering	Collect enough evidence for proportionate review	Structured intake, risk tier, accountable owners, required reviews
Pilot enablement	Test value and controls before production commitment	Pilot charter, baseline, data inventory, evaluation plan, issue log, closure record
Shared engineering patterns	Avoid duplicated or insecure implementations	Approved architectures, connector patterns, prompt/eval libraries, APIs, templates
AI asset and use-case inventory	Maintain enterprise visibility and accountability	Owners, data, models, agents, prompts, controls, status, metrics, revalidation dates
Portfolio learning and assurance	Identify value, failure modes, duplication, and control gaps	Quarterly reviews, sanitized case studies, retirement decisions, audit evidence

6. Pilot lifecycle and decision gates




Figure 3. Recommended progression from learning to operational scale.

Gate	Minimum question	Required evidence
Purpose and owner	Is the problem important, and who remains accountable?	Problem statement, process owner, affected users, alternatives considered
Data authority	May the proposed information be used for this purpose?	Data inventory, classification, owner authorization, contractual/regulatory constraints
Exact platform and configuration	Is the licence, model, feature, region, connector, and retention setting approved?	Configuration record, contract/risk reference, architecture and data flow
Human accountability	Who verifies outputs and prevents inappropriate reliance?	Review checklist, authoritative sources, prohibited uses, escalation path
Evaluation	What baseline, quality measures, risk tests, and stopping conditions apply?	Representative test set, success thresholds, issue log, red-team/boundary tests where needed
Operational readiness	Can support, security, records, continuity, and ownership be sustained?	Support model, monitoring, incident plan, access lifecycle, change process
Closure or scale	What happens to access, files, agents, indexes, records, and learning?	Closure report, deletion/retention evidence, decision and conditions for scale

7. Workforce enablement and adoption

Corporate examples show that adoption is not produced by licensing alone. Walmart used a playground and familiar productivity tasks; Moderna set a proficiency objective and tested tools with users; Morgan Stanley relied on domain experts to evaluate outputs; JPMorganChase pairs central tools with embedded business teams. These patterns support role-based learning integrated with real work rather than generic prompt instruction alone.

Audience	Essential learning and responsibilities
All employees	Recognize approved environments; classify and minimize information; verify outputs; protect confidential and personal information; report incidents.

Audience	Essential learning and responsibilities
Managers and sponsors	Select bounded workflows; address workload and role changes; define accountability; set measures; decide whether to stop, revise, or scale.
Business process experts	Document authoritative sources, exceptions, failure modes, and quality criteria; participate in evaluations.
Developers and technical teams	Secure architecture, secrets, APIs, retrieval security, prompt injection, testing, monitoring, rollback, and change control.
Data owners and privacy/legal/risk teams	Purpose compatibility, authorization, proportionality, retention, third-party terms, cross-border handling, and consequential-use restrictions.
AI champions / facilitators	Translate policy into practical guidance, collect recurring questions, support peer learning, and route issues without becoming informal approval authorities.

8. Measures of service quality and value

Metrics should not be limited to licence activation or message volume. The service should measure whether employees make better decisions, pilots produce reliable evidence, and controls remain effective.

Outcome	Illustrative measures
Accessible service pathway	Time to initial triage; percentage routed correctly; backlog age; user comprehension.
Appropriate platform and data use	Reduction in personal-account or unreviewed-tool use; classification accuracy; policy exception trends.
Pilot evidence quality	Baseline completeness; representative evaluation coverage; quality/error/rework; documented stop/revise/scale decisions.
Business value	Cycle time, quality, customer or employee experience, reviewer effort, avoided work, revenue or risk outcome where valid.
Control effectiveness	Incidents and near misses; DLP findings; permission defects; unresolved vulnerabilities; revalidation timeliness.
Reusable organizational learning	Approved patterns, evaluation datasets, sanitized failure cases, duplicate projects avoided, retired weak use cases.
Workforce and fairness	Access across roles/regions, training completion and demonstrated proficiency, workload transfer, accessibility, employee trust.

9. Institution-agnostic enterprise roadmap

Period	Priority actions
0-60 days	Confirm executive sponsor and decision rights; inventory existing tools and shadow use; define approved data ceilings; establish an interim safe-use standard; design intake and risk tiers; select a protected learning environment.
2-4 months	Launch role-based training, office hours, and champions; create pilot kit; establish cross-functional triage; select reversible pilots; create the initial use-case and AI-asset inventory.
4-9 months	Evaluate pilots; publish approved patterns and sanitized failure cases; implement persistent-assistant and connector checklists; establish evaluation libraries, monitoring, and vendor-change review.

Period	Priority actions
9-18 months	Standardize successful patterns; expand platform capacity selectively; integrate DLP, records, and assurance workflows; audit repository permissions and access recertification; retire weak or duplicative assets.
18-24 months	Conduct an independent program review; revise funding and staffing; rationalize platforms; mature service catalogue, portfolio governance, board reporting, and continuous assurance.

10. Actionable recommendations

1. Create one enterprise front door for AI questions, learning access, pilot proposals, procurement, and production deployments.
2. Publish a tool-by-data-by-interaction matrix. A tool list alone does not capture differences among protected chat, uploaded files, persistent assistants, connectors, APIs, and automated actions.
3. Provide an enterprise-managed learning environment early to reduce unmanaged consumer-tool use, but begin exercises with public or synthetic information.
4. Use a hub-and-spoke model: central governance and platform teams set common controls; business units own workflow purpose, authorized data, adoption, review, and outcomes.
5. Require a named owner, data inventory, human-review procedure, evaluation plan, issue log, and closure date for every persistent assistant or pilot.
6. Make repository permission review and security-trimming tests prerequisites for retrieval-connected AI.
7. Create representative evaluation datasets and regression tests for high-value or regulated use cases, following the pattern publicly described by Morgan Stanley.
8. Implement governance as a workflow and product capability: structured intake, risk tiering, evidence capture, status, approvals, monitoring, and revalidation.
9. Inventory AI assets across vendors, including models, prompts, agents, connectors, APIs, datasets, evaluations, owners, and production status.
10. Use adoption metrics to diagnose learning and support needs, not as the sole measure of value. Pair usage with quality, risk, workload, and business-outcome evidence.
11. Publish sanitized stopped-pilot and failure-case summaries. Terminating a weak use case should be treated as useful governance evidence.
12. Require revalidation when models, vendors, features, connectors, data sources, permissions, regulations, or business purposes materially change.
13. For agentic systems, introduce transaction limits, least-privilege service identities, human approval for consequential actions, rollback, and continuous monitoring before scale.

14. Plan decommissioning from the beginning: access removal, file and index disposition, records retention, prompt/agent ownership transfer, and user communication.

11. Limitations

Public corporate materials are incomplete and may emphasize successful initiatives. Vendor case studies may be promotional and are not independent audits. Companies generally do not publish detailed internal policies, architecture diagrams, incident histories, contractual terms, evaluation datasets, or rejected use cases. The report therefore identifies publicly supported patterns rather than claiming to reconstruct each company's full operating model.

Recommendations must be adapted to sector-specific regulation, employment law, customer obligations, intellectual-property rules, records and litigation requirements, cybersecurity architecture, geographic operations, corporate risk appetite, and collective bargaining obligations where applicable.

Sources and evidence notes

[C-1] JPMorganChase, 2024 Annual Report, letter from Jennifer A. Piepszak. Supports launch of LLM Suite to more than 200,000 employees, controlled environment, and shared capabilities.

<https://www.jpmorganchase.com/ir/annual-report/2024/ar-ceo-letter-jennifer-piepszak>

[C-2] JPMorganChase, Data & AI ecosystem careers page. Supports centralized CDAO strategy/governance/tooling and embedded business teams using approved tools.

<https://www.jpmorganchase.com/careers/explore-opportunities/programs/data-analytics-opportunities>

[C-3] Morgan Stanley, Firmwide Artificial Intelligence. Supports a firmwide AI team, human-centric design, data/system protection, controls, and oversight.

<https://www.morganstanley.com/about-us/technology/artificial-intelligence-firmwide-team>

[C-4] Morgan Stanley, AI @ Morgan Stanley Debrief launch. Supports sequence from Assistant to Debrief, reported 98% adviser-team adoption, employee policy obligations, output limitations, and verification. <https://www.morganstanley.com/press-releases/ai-at-morgan-stanley-debrief-launch>

[C-5] Walmart, Expanding My Assistant. Supports 60-day initial development, launch to 50,000 U.S. campus associates, international expansion, use cases, and feedback.

<https://corporate.walmart.com/news/2024/01/09/walmarts-expanding-one-of-a-kind-associate-genai-tool-to-11-countries-in-2024>

[C-6] Walmart Global Tech, Empowering developers with AI. Supports GenAI Playground as an internal safe experimentation environment and the subsequent My Assistant launch.

https://public.walmart.com/content/walmart-global-tech/en_us/blog/post/empowering-developers-with-ai.html

[C-7] Unilever, AI Governance Product Lead role description. Supports governance-platform capabilities: use-case intake, risk triage, assessment, assurance evidence, lifecycle monitoring, metrics,

and cross-functional collaboration. <https://careers.unilever.com/es/trabajo/bengaluru/ai-governance-product-lead/34155/97609301824>

[C-8] Unilever, People Experience Innovation Funnel Manager role description. Supports structured progression from discovery to proof of concept, pilot, and global scale-up.

<https://careers.unilever.com/es/trabajo/kattowitz/people-experience-innovation-funnel-manager/34155/92077389616>

[C-9] Unilever, Head of Artificial Intelligence role description. Supports AI council, governance framework, oversight, security, data protection, data governance, training, and a centre of excellence.

<https://careers.unilever.com/job/kingston-upon-thames/ice-cream-head-of-artificial-intelligence/34155/76995238640>

[C-10] IBM documentation, Governing assets with watsonx.governance. Supports end-to-end use-case tracking, factsheets, roles, approval workflow, evaluation, monitoring, and third-party model support.

<https://www.ibm.com/docs/en/watsonx/saas?topic=governing-ai>

[C-11] IBM documentation, watsonx.governance use case. Supports tracking, evaluation, and monitoring of models and prompt templates.

<https://www.ibm.com/docs/en/watsonx/saas?topic=cases-watsonxgovernance-use-case>

[C-12] IBM product page, watsonx.governance. Supports IBM-reported internal reuse approvals and data-clearance processing improvement; treated as a vendor claim.

<https://www.ibm.com/products/watsonx-governance>

[V-1] OpenAI customer case study, Morgan Stanley. Supports vendor-reported evaluation practices, regression tests, zero-data-retention arrangement, and adoption outcomes.

<https://openai.com/index/morgan-stanley/>

[V-2] OpenAI customer case study, Moderna. Supports vendor-reported mChat adoption, user testing, workforce proficiency goal, employee-created GPTs, and selected use cases.

<https://openai.com/index/moderna/>

[G-1] OpenAI, Introducing ChatGPT Enterprise. Supports general enterprise features including non-training terms, encryption, SSO, domain verification, and administrative controls; product terms should be validated against current contracts.

<https://openai.com/index/introducing-chatgpt-enterprise/>

[G-2] OpenAI, New compliance and administrative tools for ChatGPT Enterprise. Supports compliance API, audit and authentication logs, and DLP/e-discovery integration direction.

<https://openai.com/index/new-tools-for-chatgpt-enterprise/>