

Launching an Institutional AI Advisory, Literacy, and Pilot Service

Institution-agnostic recommendations for higher-education institutions, with emphasis on
protected institutional information

Research-based recommendations / August 2026

Scope note: This report distinguishes published facts, comparative observations, and recommendations. It is not legal advice or a substitute for institutional privacy, security, procurement, records, accessibility, labour, research ethics, or legal review.

Executive summary

Higher-education institutions are establishing AI services through several models: enterprise licences, institution-managed multi-model environments, secure test environments, central advisory teams, role-based training, and structured pilot support. Across the reviewed institutions, no single control is treated as sufficient. The recurring model combines contractual protections, institutional identity, data classification, authorization, secure configuration, minimum-necessary information, human accountability, monitoring, and periodic reassessment. [P-1]–[P-10]

Recommended launch principle

Build a service pathway, not merely a tool catalogue. Users should be able to move from a question to a clear decision: learn safely, use an approved protected service, enter a reviewed pilot, or stop/refer the proposal. Governance should be proportionate to persistence, integration, data sensitivity, and potential impact.

1. Scope and method

This report synthesizes official public materials from selected universities and Canadian public authorities. The institutions represent different operating models: Harvard’s sandbox and school-level access controls; Michigan’s university-built service suite; Toronto’s response-team and AI Kitchen concept; Stanford’s multi-model playground, connectors, and enterprise education licences; Berkeley and Illinois tool-by-data-classification guidance; and Waterloo’s coordinated approval and pilot-oriented Lab. Public information is uneven, so this report does not rank institutions or claim that published pages fully describe internal practice.

Evidence category	Use
Published fact	A statement directly supported by an official source.
Cross-institutional observation	A pattern present in multiple official sources.
Recommendation	A proposed general practice that must be adapted to local law, policy, architecture, and governance.

2. Comparative findings

Institution/model	Published approach	Operational implication
Harvard	Institutional AI Sandbox; approved data ceiling; some school-level training and access request requirements. [P-1][P-2]	Protected environments can be paired with gated access and local controls.
Michigan	University-provided GPT, Maizey, and Toolkit services; contracts/data agreements required for sensitive data. [P-3][P-4]	A university-owned service layer can separate general chat, data-grounded experiences, and advanced development.
Toronto	AI Response Team plus AI Kitchen; enterprise tools assigned data ceilings. [P-5][P-6]	Advisory capacity and secure experimentation should develop together.
Stanford	Institution-managed multi-model playground, SSO, risk-category approvals, connectors, and enterprise education licences. [P-7][P-8][P-9]	The same institution may approve different services and connectors for different data categories.
Berkeley	Licensed tools tied to protection levels; unit-provided tools must publish limitations. [P-10]	Contractual protections apply only when users enter through the institutional service.
Illinois	Public matrix of tools and approved data classes; high-risk data excluded from AI tools. [P-11]	A simple public matrix can translate policy into operational choices.
Waterloo	Approved/limited/problematic/unreviewed categories, coordinated review, system-deployment requirements, and a pilot/advisory Lab. [P-12][P-13]	A service can build on existing governance rather than create a parallel approval authority.

3. Control architecture




Figure 1. Institutional controls should operate as a coordinated set. Vendor non-training terms address only part of the risk.

3.1 Contractual and vendor controls

- Define use of prompts, files, outputs, feedback, telemetry, and metadata for training or service improvement.
- Document retention, deletion, subprocessors, data location, breach notification, audit, intellectual property, indemnity, termination, and export.
- Confirm whether commitments vary by model, feature, connector, API, region, licence tier, or administrator setting.

- Require notice and reassessment for material product or data-handling changes.

3.2 Identity and access

- Use institutional single sign-on and multifactor authentication.
- Provision and remove access according to affiliation and role.
- Use groups or workspaces for pilots; avoid unmanaged personal accounts.
- Review administrator access, support access, delegated access, and conversation-sharing features.

3.3 Data authority and classification

- Classify information before selecting a service.
- Separate platform approval from authority to use a particular dataset for a particular purpose.
- Identify the data owner or steward and any research, contractual, donor, library-licence, employment, student, or records constraints.
- Treat inferred personal information as personal information where applicable; public availability does not automatically authorize unrestricted reuse. [G-1]

3.4 Technical configuration

- Map data flows, storage, logs, model providers, plugins, connectors, and write-back actions.
- Minimize prompts and datasets; prefer synthetic or de-identified information where effective.
- Set retention, history, sharing, export, model, region, and connector settings centrally where possible.
- Protect API keys, secrets, service identities, and administrative interfaces.

3.5 Human and organizational controls

- Define what outputs require verification and what authoritative sources must be consulted.
- Prevent consequential decisions from relying solely on model output.
- Provide disclosure, explanation, appeal, or correction pathways where people may be affected.
- Assess accessibility, bias, workload redistribution, academic freedom, labour, and professional-accountability implications.

4. Classify the interaction pattern, not only the tool

Tier 1 — Learning

Public or synthetic information; no system integration | Open workshops and self-service guidance

Tier 2 — Protected productivity

Approved enterprise tool; authorized internal information | Standard controls, role-based training, local owner

Tier 3 — Data-connected pilot

Persistent files, agents, RAG, connectors, or APIs | Formal intake, data inventory, technical review, closure plan

Tier 4 — High-impact or highly sensitive

Decisions about people, regulated data, automated actions | Specialist governance; possible exclusion or enhanced approval

Figure 2. Suggested service tiers based on information sensitivity, persistence, integration, and impact.

Interaction pattern	Examples	Why risk changes
Ephemeral protected chat	Drafting, summarizing, brainstorming in an approved account.	Limited integration, but authorization, retention, and output handling still matter.
Uploaded-file analysis	One or more documents attached to a conversation.	Files may persist, contain hidden metadata, or create records and copyright obligations.
Persistent agent	Custom instructions, knowledge files, shared assistant, memory.	Content and access can become stale; ownership and decommissioning are required.
Retrieval over repositories	RAG over document stores, LMS, research repositories, or business systems.	Search can amplify pre-existing oversharing and cross-document inference.
API workflow	Batch processing, orchestration, custom interface, database integration.	Introduces logs, secrets, error handling, subprocessors, and software lifecycle risks.
Agentic or automated action	Email, ticket, record, purchase, scheduling, or decision action.	Errors can propagate directly into institutional systems and affect people.
Model customization	Fine-tuning, embeddings, persistent vector stores, local models.	Creates new data assets, retention obligations, evaluation needs, and attack surfaces.

5. Recommended service model

Service component	Purpose	Typical outputs
Advisory front door	Clarify the problem, route requests, reduce hidden experimentation.	Triage decision, consultation, referral, use-case framing.
Literacy and training	Develop practical capability for different roles.	Workshops, office hours, scenario exercises, role guides, train-the-trainer.
Protected experimentation	Allow safe learning with synthetic, public, or authorized information.	Sandbox/playground access, sample datasets, test workspaces.
Pilot enablement	Test value and controls before operational commitment.	Charter, baseline, data inventory, review plan, issue log, closure report.
Shared capabilities	Avoid duplicate procurement and inconsistent configurations.	Approved platforms, APIs, patterns, templates, connector standards.
Portfolio learning	Identify demand, duplication, failure modes, and capability gaps.	Use-case register, quarterly review, pattern library, retirement decisions.

6. Governance and decision rights

A central AI team should coordinate but should not become the sole owner of every risk. A practical responsibility structure is:

Decision	Accountable role (generic)	Contributors
Is the problem worth addressing?	Business, academic, research, or service owner	Users, process experts, affected groups.
May the data be used for this purpose?	Data owner/steward or delegated authority	Privacy, legal, research ethics, records.
Is the platform/configuration acceptable?	Technology/security authority	Architecture, procurement, vendor management.
Is the pilot method adequate?	Pilot sponsor	AI service, evaluation, accessibility, risk specialists.
May the service become operational?	Service/system owner	All relevant reviewers; affected operational units.
Should the service continue after changes?	Service owner with governance body	Vendor management, security, privacy, data owners, users.

7. Pilot lifecycle and gates

Decision gates: data authority • security/privacy • risk proportionality • human accountability



Every pilot should produce reusable evidence, not only a local demonstration.

Figure 3. Recommended lifecycle. Review depth should be proportionate, but no pilot should begin without ownership, permitted data, and a closure plan.

Gate	Minimum question	Evidence
Purpose	Is AI necessary and proportionate to the stated problem?	Problem statement, alternatives, expected benefit.
Authority	Who owns the data and authorizes this use?	Owner/steward confirmation and classification.
Platform	Is the exact service, licence, model, region, feature, and connector approved?	Configuration record and contract/risk reference.
Data minimization	Can synthetic, public, de-identified, or reduced data answer the question?	Data inventory and minimization decision.
Human accountability	Who verifies output and prevents inappropriate reliance?	Review procedure, escalation and prohibited uses.
Evaluation	What baseline and stop conditions will be used?	Measures, sample, timeframe, issue log.
Closure	How will access, data, agents, indexes, and records be handled?	Deletion/retention and transition plan.

8. Repository-connected AI: specific controls

Repository-connected services deserve a separate control track. A model that respects source permissions can still expose information inappropriately if the source repository is overshared, if group membership is stale, or if retrieval combines fragments in ways users did not anticipate.

- Identify each indexed source, owner, classification, purpose, and authoritative status.
- Review inherited permissions, anonymous links, broad groups, external guests, stale accounts, and documents shared beyond operational need.
- Confirm whether security trimming applies at retrieval time and how changes propagate to indexes and caches.
- Test adversarial and boundary queries, including attempts to retrieve other users’ information or hidden instructions.
- Display source links or citations where feasible so users can verify context and authority.
- Define how deleted or reclassified source content is removed from indexes, embeddings, caches, logs, and backups.
- Restrict agent sharing and connectors by group; establish periodic access recertification.

9. Training framework

Audience	Essential learning
All users	Approved account recognition; classification; minimization; verification; copyright; disclosure; incident reporting.
Pilot participants	Use-case boundaries; human-review procedure; permitted data; issue logging; stopping conditions.
Managers and sponsors	Accountability, workload, labour, accessibility, evaluation, scaling and retirement decisions.
Technical teams	Architecture, secrets, APIs, logging, retrieval security, prompt injection, evaluation, change control.
Data owners/stewards	Purpose compatibility, authorization, proportionality, retention, de-identification, re-identification risk.
Procurement/vendor managers	Terms, subprocessors, training, retention, model/feature changes, audit, exit and data return.

10. Measures of service quality

Outcome	Measures
Understandable pathway	Time to initial decision; percentage routed correctly; user comprehension.
Appropriate use	Tool/data selection accuracy; reduction in personal-account or unreviewed-tool use.
Pilot evidence	Baseline completeness; quality/error/rework measures; documented stop/revise/scale decisions.
Risk responsiveness	Near-miss reporting; issue closure; connector permission findings; revalidation timeliness.
Reusable learning	Published patterns, checklists, sanitized failure cases, duplicate projects avoided.
Institutional fairness	Access across roles and units; accommodation; support burden; workload and employment effects.

11. Implementation roadmap

Period	Priority actions
0–60 days	Confirm governance; inventory tools and contracts; define data ceilings; select service tiers; design intake; publish interim safe-use guidance.
2–4 months	Launch training and office hours; create pilot kit; establish triage meeting; choose initial reversible pilots; implement use-case register.
4–9 months	Evaluate pilots; publish patterns and failure cases; implement connector checklist; refine role training; establish change/revalidation monitoring.
9–18 months	Standardize successful patterns; retire weak ones; expand capacity selectively; audit access, records, equity, and control effectiveness.
18–24 months	Independent review; funding and staffing decision; platform rationalization; mature service catalogue and governance cadence.

12. Actionable recommendations

1. Publish a tool-by-data-by-interaction matrix rather than a tool list alone.
2. Create a single front door that routes learning, tool questions, pilots, procurement, and system deployments.
3. Use synthetic or public datasets for initial training and feasibility work; introduce protected information only when necessary and authorized.
4. Require a pilot charter, data inventory, human-review plan, baseline, issue log, and closure record.
5. Establish a separate review checklist for persistent agents, retrieval connectors, APIs, and automated actions.
6. Make repository permission review a prerequisite for data-connected retrieval.
7. Define service-level targets for triage and specialist referral so governance does not become an unbounded queue.
8. Publish sanitized failure cases and stopped pilots as evidence of effective governance.
9. Monitor material vendor, model, feature, connector, and contract changes and trigger proportionate revalidation.
10. Evaluate the service on safe decisions, quality, workload, accessibility, and institutional learning—not licence consumption alone.

13. Limitations and adaptation requirements

The framework must be adapted to the institution’s jurisdiction, public- or private-sector status, collective agreements, data-classification standard, records law, research ethics, health and student-information requirements, accessibility obligations, procurement rules, and architecture. Public institutional pages provide useful models but do not disclose all implementation details. Institutions should validate vendor claims against executed contracts, administrator settings, technical tests, and current documentation.

Sources and evidence notes

[P-1] Harvard AI Sandbox. <https://www.huit.harvard.edu/ai-sandbox> — Supports institution-managed access and approved information-classification ceilings.

[P-2] Harvard Medical School generative AI guidance. <https://it.hms.harvard.edu/about/policies-and-guidelines/generative-ai> — Supports school-level service matrices, training, and access requirements.

[P-3] U-M Generative AI service suite. <https://genai.umich.edu/> — Supports differentiated institutional services for general chat, data-grounded experiences, and advanced control.

[P-4] Artificial Intelligence and U-M Data. <https://safecomputing.umich.edu/protect-the-u/safely-use-sensitive-data/AI-and-UM-Data> — Supports contracts/data agreements for sensitive information and human review of generated code.

[P-5] University of Toronto AI Task Force report. <https://ai.utoronto.ca/u-of-t-ai-task-force/report/> — Supports the AI Response Team and AI Kitchen concepts.

[P-6] University of Toronto ChatGPT Edu guide. <https://teaching.utoronto.ca/tool-guides/chatgpt-edu/> — Supports a specific enterprise tool data ceiling and non-training statement.

[P-7] Stanford AI Playground. <https://uit.stanford.edu/service/aiplayground> — Supports an SSO-protected, institution-managed multi-model environment.

[P-8] Stanford AI Playground FAQs. <https://uit.stanford.edu/aiplayground/faqs> — Supports vendor retention distinctions and local restrictions.

[P-9] Stanford approved connectors for AI services. <https://uit.stanford.edu/ai/connectors> — Supports connector-specific availability and risk categories.

[P-10] Berkeley appropriate use of generative AI tools. <https://oercs.berkeley.edu/appropriate-use-generative-ai-tools> — Supports licensed-account protections and protection-level limitations.

[P-11] Illinois AI Apps. <https://genai.illinois.edu/ai-apps/> — Supports a public tool-by-data-classification matrix.

[P-12] Waterloo Approval & Support. <https://uwaterloo.ca/information-systems-technology/about/policies-standards-and-guidelines/responsible-use-ai-tools-university-data/approval-support> — Supports coordinated review, contracting, approval, and revalidation.

[P-13] Waterloo AI Discovery Lab. <https://aidiscoverylab.uwaterloo.ca/> — Supports advisory, literacy, pilot, and shared-capability service functions.

[G-1] Canadian privacy regulators: principles for generative AI. https://www.priv.gc.ca/en/privacy-topics/technology/artificial-intelligence/gd_principles_ai/ — Supports legal authority, necessity/proportionality, minimization, retention, transparency, and human oversight.

[G-2] Canada.ca privacy and security for AI help applications. <https://design.canada.ca/guidance/ai/privacy-security.html> — Supports early privacy engagement, input controls, retention/deletion, logging, and monitoring.